

Building and Training a Machine Learning Model, Containerizing, and Deploying a Machine Learning Model Using Flask, Docker Images, and Kubernetes Services

This project demonstrates the end-to-end process of building, training, and deploying a machine learning model within a robust DevOps environment. The workflow involves:

1. Building and Training the Machine Learning Model:

- A simple machine learning model is created using Python and libraries like scikit-learn.
- The model is trained with sample data, serialized into a model.pkl file, and saved for deployment.

2. Creating a Flask API for Model Serving:

- A Flask application is developed to serve the trained model, providing endpoints for making predictions.
- The API processes input data, loads the serialized model, and returns predictions in real-time.

3. Containerizing the Application Using Docker:

- A Docker image is created for the Flask application, bundling the app code, dependencies, and the trained model.
- Docker commands are used to build, test, and run the container locally to ensure functionality.

4. Deploying the Containerized Application Using Kubernetes:

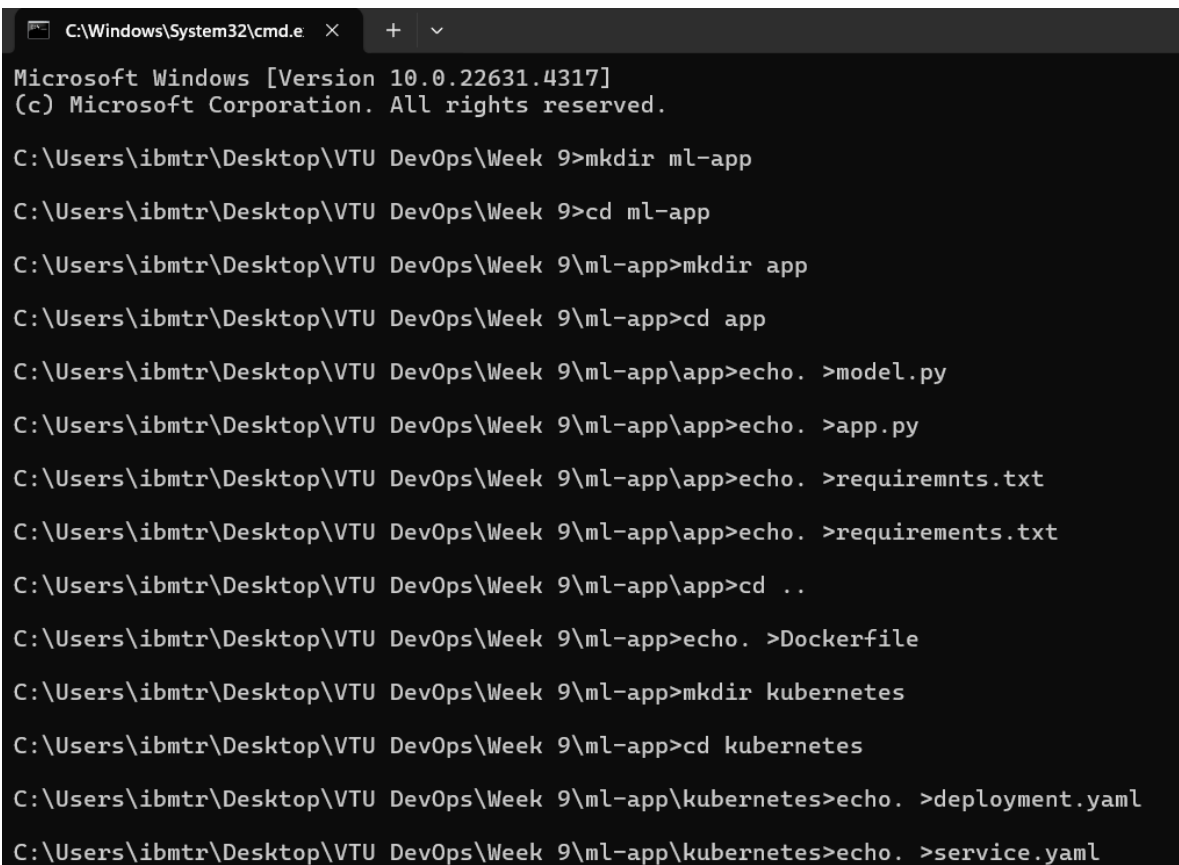
- Minikube is used to simulate a Kubernetes cluster locally.
- The Docker image is deployed as a Kubernetes Pod, with a Service created to expose the application.
- kubectl commands manage the deployment, scaling, and health of the Pods.

Step 1: Create the ML Model

1. Directory structure:

ml-app/

```
|— app/
|   |— model.py
|   |— app.py
|   └— requirements.txt
|— Dockerfile
|— kubernetes/
|   |— deployment.yaml
|   └— service.yaml
└— README.md
```



```
C:\Windows\System32\cmd.e  X  +  v
Microsoft Windows [Version 10.0.22631.4317]
(c) Microsoft Corporation. All rights reserved.

C:\Users\ibmtr\Desktop\VTU DevOps\Week 9>mkdir ml-app
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9>cd ml-app
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>mkdir app
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>cd app
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app\app>echo. >model.py
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app\app>echo. >app.py
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app\app>echo. >requiremnts.txt
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app\app>echo. >requirements.txt
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app\app>cd ..
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>echo. >Dockerfile
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>mkdir kubernetes
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>cd kubernetes
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app\kubernetes>echo. >deployment.yaml
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app\kubernetes>echo. >service.yaml
```

2. ML model code (model.py):

```
import pickle

from sklearn.datasets import load_iris

from sklearn.ensemble import RandomForestClassifier

# Train and save a model

def train_model():

    try:

        # Load dataset

        data = load_iris()

        X, y = data.data, data.target

        # Train a simple model

        model = RandomForestClassifier()

        model.fit(X, y)

        # Save the model to a file

        with open("model.pkl", "wb") as file:

            pickle.dump(model, file)

        print("Model created and saved as 'model.pkl'")

    except Exception as e:

        print(f"Error: {e}")

if __name__ == "__main__":

    train_model()
```

3. **Flask app code (app.py):**

```
import pickle

from flask import Flask, request, jsonify

app = Flask(__name__)

# Load the trained model

with open("model.pkl", "rb") as file:

    model = pickle.load(file)

@app.route("/")

def home():

    return "Welcome to the ML App!"

@app.route("/predict", methods=["POST"])

def predict():

    data = request.json

    prediction = model.predict([data["features"]])

    return jsonify({"prediction": prediction.tolist()})

if __name__ == "__main__":

    app.run(host="0.0.0.0", port=5000)
```

4. **Requirements file (requirements.txt):**

Flask==2.3.3

scikit-learn==1.3.2

5. **Train the model:** Run the model.py script to generate the model.pkl file:

python app/model.py

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>python app/model.py
Model created and saved as 'model.pkl'
```

Desktop > VTU DevOps > Week 9 > ml-app >				
Sort View ...				
Name	Date modified	Type	Size	
app	19-11-2024 17:19	File folder		
kubernetes	19-11-2024 17:20	File folder		
Dockerfile	19-11-2024 17:27	File	1 KB	
model.pkl	19-11-2024 17:33	PKL File	185 KB	

Step 2: Create a Docker Image

1. Dockerfile:

FROM python:3.9-slim

WORKDIR /app

Copy the application and the model file into the Docker image

COPY app/ /app/

COPY model.pkl /app/

RUN pip install -r requirements.txt

EXPOSE 5000

CMD ["python", "app.py"]

Step 3: Verify Image Availability

Ensure Kubernetes can access the image:

docker images

2. Build the Docker image:

docker build -t ml-app .

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>docker build -t ml-app .
[+] Building 4.3s (11/11) FINISHED
=> [internal] load build definition from Dockerfile                                docker:desktop-linux 0.2s
=> => transferring dockerfile: 263B                                              0.1s
=> [internal] load metadata for docker.io/library/python:3.9-slim                3.5s
=> [auth] library/python:pull token for registry-1.docker.io                    0.0s
=> [internal] load .dockerignore                                                  0.0s
=> => transferring context: 2B                                                    0.0s
=> [1/5] FROM docker.io/library/python:3.9-slim@sha256:6250eb7983c08b3cf5a7db9309f8630d3ca03dd152158fa37a3f8daaf 0.0s
=> [internal] load build context                                                  0.0s
=> => transferring context: 156B                                                 0.0s
=> CACHED [2/5] WORKDIR /app                                                      0.0s
=> CACHED [3/5] COPY app/ /app/                                                    0.0s
=> CACHED [4/5] COPY model.pkl /app/                                              0.0s
=> CACHED [5/5] RUN pip install -r requirements.txt                              0.0s
=> exporting to image                                                            0.0s
=> => exporting layers                                                            0.0s
=> => writing image sha256:b03c2e82a8aa621c8baf41a6b7feb4a15aaa2f814ff1c8b3b4e92f8c7c998aa8 0.0s
=> => naming to docker.io/library/ml-app                                         0.0s

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/hqjai3kbmyhasej12wj50q1yo

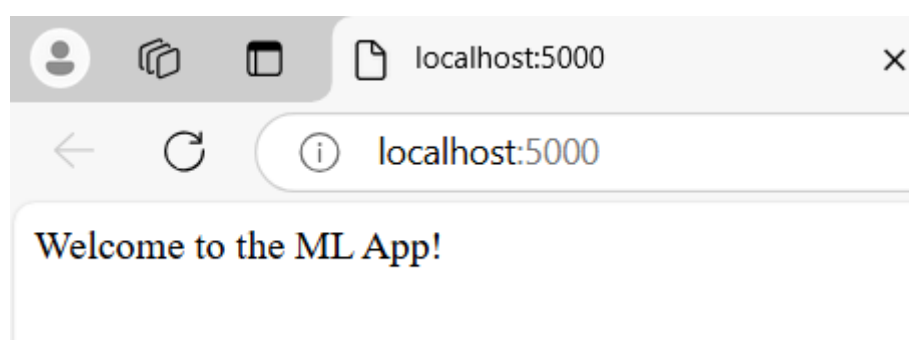
What's next:
  View a summary of image vulnerabilities and recommendations → docker scout quickview

C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>
```

3. Test the Docker container locally:

docker run -p 5000:5000 ml-app

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>docker run -p 5000:5000 ml-app
* Serving Flask app 'app'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Running on http://172.17.0.2:5000
Press CTRL+C to quit
172.17.0.1 - - [19/Nov/2024 12:07:09] "GET / HTTP/1.1" 200 -
172.17.0.1 - - [19/Nov/2024 12:07:09] "GET / HTTP/1.1" 200 -
172.17.0.1 - - [19/Nov/2024 12:07:11] "GET /favicon.ico HTTP/1.1" 404 -
172.17.0.1 - - [19/Nov/2024 12:07:11] "GET /favicon.ico HTTP/1.1" 404 -
```




If you're using a remote image, ensure it is pushed to a container registry like Docker Hub:

`docker tag ml-app:latest <your-dockerhub-username>/ml-app:latest`

`docker push <your-dockerhub-username>/ml-app:latest`

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>docker tag ml-app:latest divyamurugan/ml-app:latest


C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>docker push divyamurugan/ml-app:latest
The push refers to repository [docker.io/divyamurugan/ml-app]
31be0e4c7d7e: Pushed
22af955e1935: Pushed
0af63ec193bc: Pushed
0cbe6fec9a4d: Pushed
aacba17e24d9: Mounted from library/python
f751ad7c65c4: Mounted from library/python
7822e749b484: Mounted from library/python
c3548211b826: Mounted from library/python
latest: digest: sha256:4b597fd1c87876d44ca339720f4fc9ce7cd7b49847f13bfa510b545be88908ef size: 1994
```



 **dockerhub** [Explore](#) [Repositories](#) [Organizations](#) [Usage](#)

divyamurugan

All Content

Create repository

Name	Size	Last Pushed 	Contains	Visibility	Scout
divyamurugan/ml-app	0 Bytes	1 minute ago	IMAGE	Public	Inactive
divyamurugan/divya	0 Bytes	6 days ago	IMAGE	Public	Inactive

1-2 of 2  

Step 3: Deploy with Minikube and kubectl

1. Start Minikube:

minikube start

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>minikube start
* minikube v1.34.0 on Microsoft Windows 11 Home Single Language 10.0.22631.4317 Build 22631.4317
* Using the docker driver based on existing profile
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.45 ...
* Restarting existing docker container for "minikube" ...
! Failing to connect to https://registry.k8s.io/ from inside the minikube container
* To pull new external images, you may need to configure a proxy: https://minikube.sigs.k8s.io/docs/reference/networking/proxy/
* Preparing Kubernetes v1.31.0 on Docker 27.2.0 ...
* Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: default-storageclass, storage-provisioner
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>minikube status
minikube
type: Control Plane
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured
```

2. Create Kubernetes manifests:

- **Deployment (deployment.yaml):**

apiVersion: apps/v1

kind: Deployment

metadata:

name: ml-app

spec:

replicas: 1

selector:

matchLabels:

app: ml-app

template:

metadata:

labels:

app: ml-app

spec:

containers:

- name: ml-app

image: divyamurugan/ml-app:latest

ports:

- containerPort: 5000

- ☐ Update the deployment.yaml to reference the correct image:

image: <your-dockerhub-username>/ml-app:latest

- **Service (service.yaml):**

apiVersion: v1

kind: Service

metadata:

name: ml-app-service

spec:

selector:

app: ml-app

ports:

- protocol: TCP

port: 80

targetPort: 5000

type: NodePort

3. Apply the manifests:

kubectl apply -f kubernetes/deployment.yaml

kubectl apply -f kubernetes/service.yaml

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>kubectl apply -f kubernetes/deployment.yaml
deployment.apps/ml-app created
```

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>kubectl apply -f kubernetes/service.yaml
service/ml-app-service created
```

Check Deployment and Service

1. Check Deployment: Verify the status of the Deployment:

kubectl get deployments

kubectl get pods

kubectl get svc

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
ml-app	1/1	1	1	4m5s

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
ml-app-645c8c678d-nv57g	1/1	Running	0	4m10s

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5m20s
ml-app-service	NodePort	10.102.124.192	<none>	80:30638/TCP	4m3s

4. Access the service: Get the service URL using Minikube:

minikube service ml-app-service

```
C:\Users\ibmtr\Desktop\VTU DevOps\Week 9\ml-app>minikube service ml-app-service
```

NAMESPACE	NAME	TARGET PORT	URL
default	ml-app-service	80	http://192.168.49.2:30638

```
* Starting tunnel for service ml-app-service.
```

NAMESPACE	NAME	TARGET PORT	URL
default	ml-app-service		http://127.0.0.1:50799

```
* Opening service default/ml-app-service in default browser...
```

```
! Because you are using a Docker driver on windows, the terminal needs to be open to run it.
```

